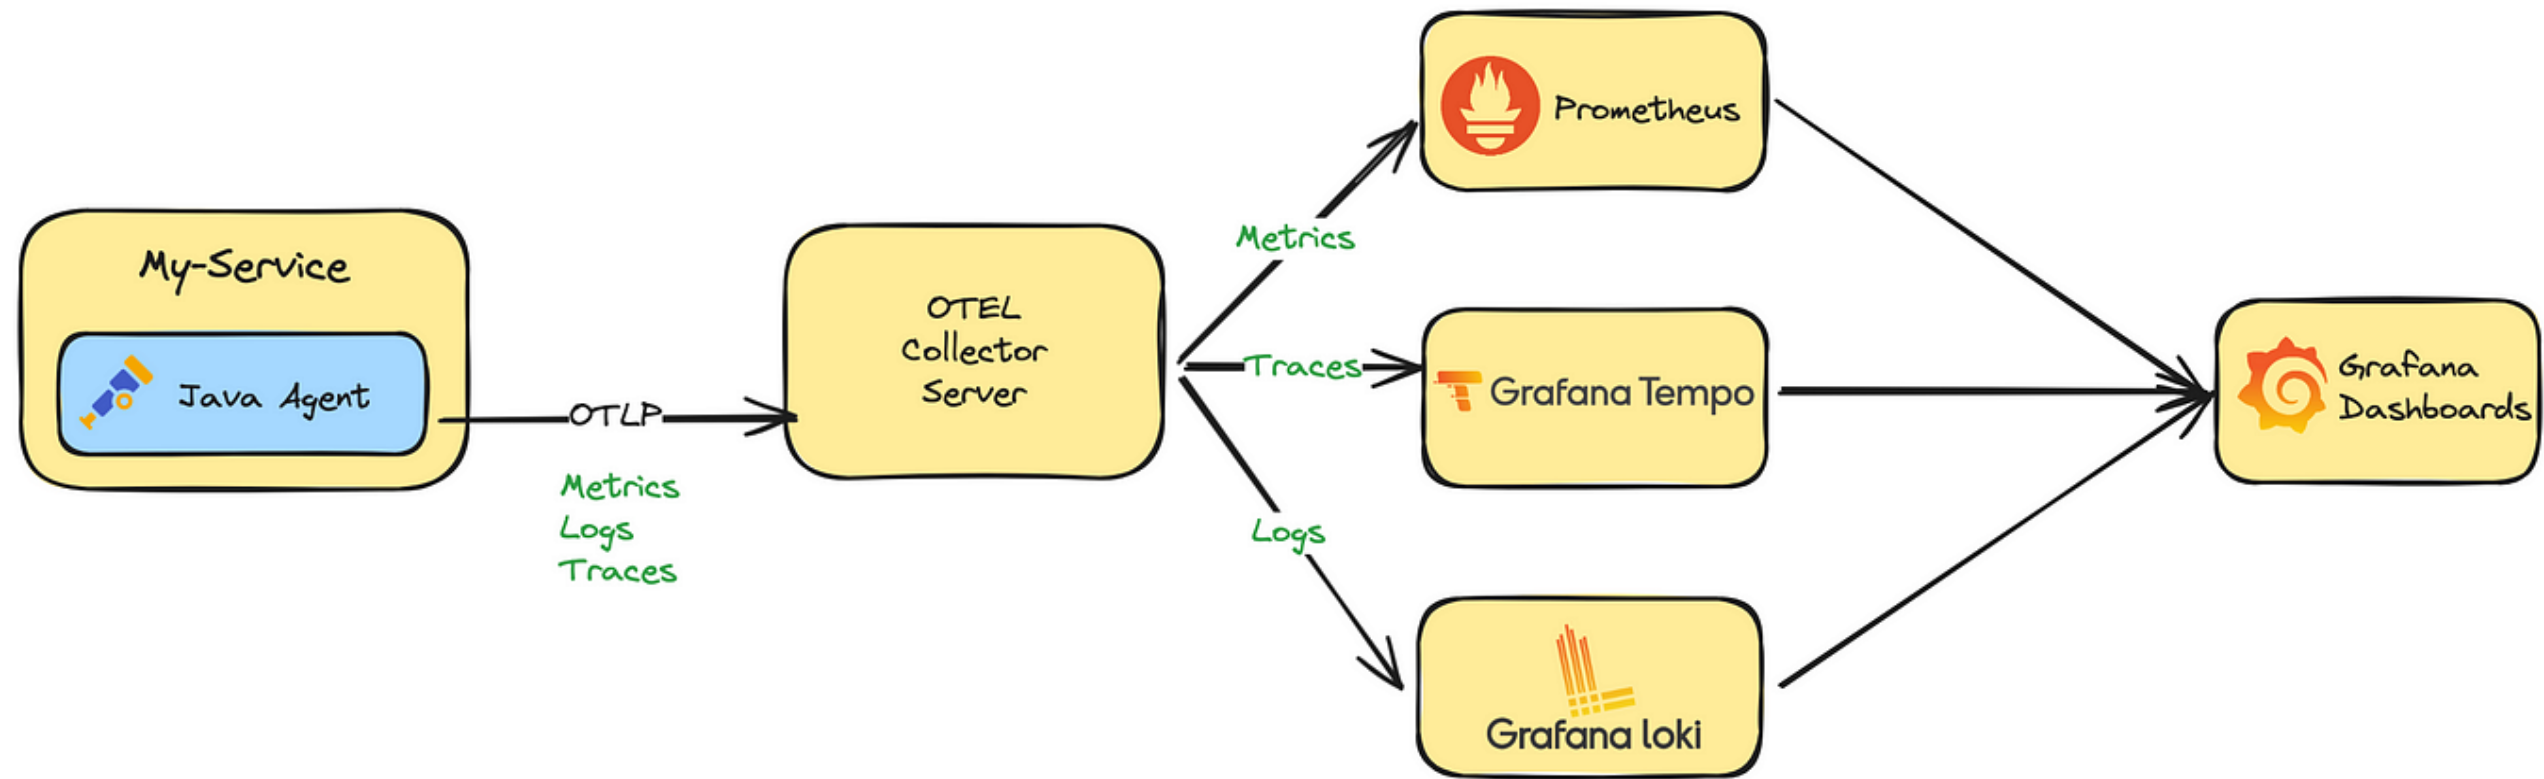
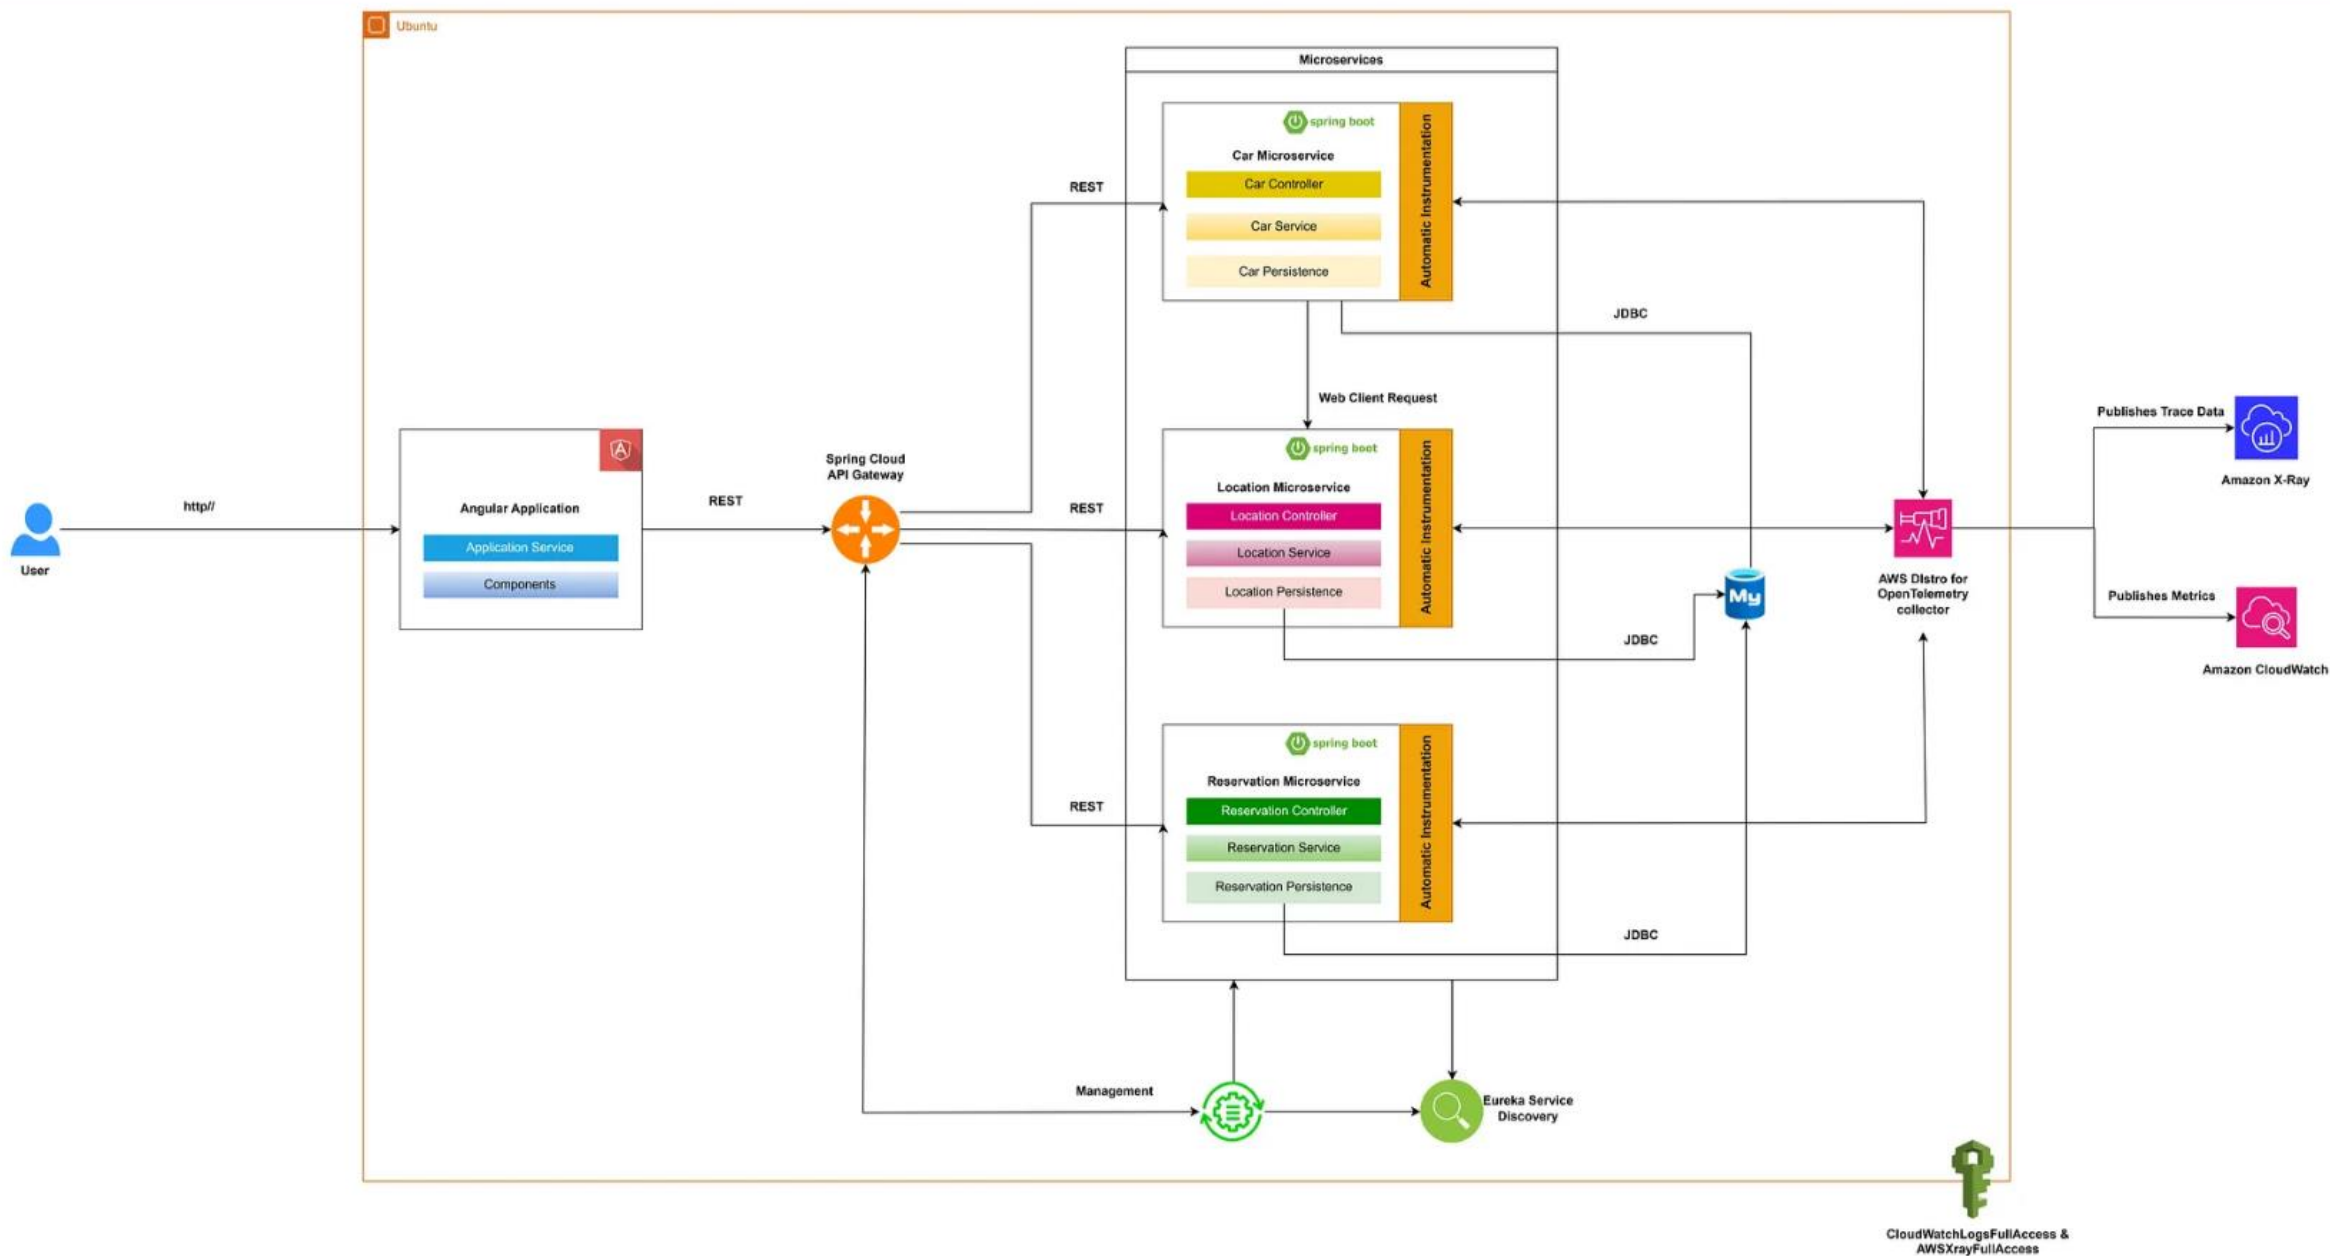


# Enterprise Essentials

## Observability - 2







## b. Docker

- **Base Image:** A lightweight Docker image ( `maven:3.8.1-openjdk-17-slim` ) containing Maven and OpenJDK 17 to build and run the Java application.
- **Dockerfile:**
  - Copies the source code into the Docker image.
  - Builds the Java application inside the container.
  - Downloads the AWS OpenTelemetry Java agent for observability.
  - Defines the entry point to run the Java application with the OpenTelemetry agent.

## 1. Architecture:

- **Microservice Architecture:** The application follows a microservice architecture using Spring Boot, making it lightweight and scalable.
- **Observability and Monitoring:** Incorporates observability through OpenTelemetry, allowing for distributed tracing and metrics collection.

## 2. OpenTelemetry Integration

- **Initialization:**
  - The `Tracer` and `Meter` instances are initialized using `GlobalOpenTelemetry`.
  - Uses configurable versions ( `otel.traces.api.version` and `otel.metrics.api.version` ) for flexibility.
- **Custom Metrics:**
  - **Counter** ( `numberOfExecutions` ): Tracks the execution count of the `/hello` endpoint.
  - **Gauge:** Monitors heap memory usage, providing insight into application resource consumption.
- **Tracing:**

## 2. Components:

- **Spring Boot Application:**
  - The application is built using Spring Boot ( `spring-boot-starter-web` ), which sets up a production-ready web service quickly.
  - It includes a web server (Tomcat by default) and an Actuator for monitoring and management endpoints.
- **Observability (OpenTelemetry Integration):**
  - Uses OpenTelemetry APIs ( `opentelemetry-api` , `opentelemetry-instrumentation-annotations` ) for capturing traces and metrics.
  - Ensures the application is instrumented to collect telemetry data for performance monitoring and troubleshooting.
- **DevTools:**
  - `spring-boot-devtools` is included for hot-reloading and easier development, enhancing developer productivity.
- **Testing:**
  - Incorporates JUnit ( `junit` ) and `spring-boot-starter-test` for unit testing and integration testing, ensuring code quality and reliability.

## 1. Receivers

```
yaml
receivers:
  otlp:
    protocols:
      grpc:
        endpoint: 0.0.0.0:5555
```

- **Receivers** are components that receive telemetry data from your application. Here, the receiver is set up to accept OTLP (OpenTelemetry Protocol) data.
- The OTLP receiver is configured to use the `grpc` protocol, listening on all network interfaces (`0.0.0.0`) at port `5555`. This means it's set up to receive telemetry data over gRPC on this port.

## 3. Exporters

```
yaml
exporters:
  prometheus:
    endpoint: collector:6666
    namespace: default
  otlp:
    endpoint: tempo:4317
  tls:
    insecure: true
```

- **Exporters** are components that send processed telemetry data to a back-end or monitoring system.
- Here, two exporters are configured:
  - **Prometheus:**
    - This exporter sends metrics to a Prometheus-compatible system.

## 2. Processors

```
yaml
processors:
  batch:
    timeout: 1s
    send_batch_size: 1024
```

- **Processors** are used to modify or transform the telemetry data. In this case, a `batch` processor is defined.
- The `batch` processor groups individual telemetry items into batches for more efficient processing and export.
  - `timeout: 1s` - The time to wait before sending a batch.
  - `send_batch_size: 1024` - The number of telemetry items to group into a batch before sending.

## 4. Service

```
yaml
service:
  pipelines:
    metrics:
      receivers: [otlp]
      processors: [batch]
      exporters: [prometheus]
    traces:
      receivers: [otlp]
      processors: [batch]
      exporters: [otlp]
  telemetry:
    logs:
      level: debug
```

- **Service** defines how telemetry data is processed through pipelines.

### 3. Exporters

Exporters send the processed telemetry data to the desired destination.

yaml

Copy code

```
exporters:
  awsemf:
    region: 'us-east-1'
    log_group_name: '/metrics/otel'
    log_stream_name: 'otel-using-java'
```

- **awsemf**: AWS CloudWatch EMF (Embedded Metric Format) exporter.
  - **region**: AWS region where metrics will be sent ( `us-east-1` ).
  - **log\_group\_name**: The name of the log group in CloudWatch ( `/metrics/otel` ).
  - **log\_stream\_name**: The name of the log stream in CloudWatch ( `otel-using-java` ).

yaml

Copy code

```
awsxray:
  region: 'us-east-1'
```

- **awsxray**: Sends traces to AWS X-Ray for analysis.
  - **region**: AWS region for X-Ray ( `us-east-1` ).

### 4. Service

Defines pipelines to process different types of telemetry data.

yaml

Copy code


```
service:
  pipelines:
    metrics:
      receivers: [otlp]
      processors: [filter, batch]
      exporters: [awsemf]
```

- **metrics pipeline**: Specifies how metrics are processed.
  - **receivers**: Uses the `otlp` receiver to collect metrics.
  - **processors**: Applies the `filter` and `batch` processors to the metrics.
  - **exporters**: Sends the processed metrics to `awsemf` (CloudWatch).

yaml

Copy code

```
traces:
  receivers: [otlp]
  processors: [batch]
  exporters: [awsxray]
```

- **traces pipeline**: Specifies how traces are processed.
  - **receivers**: Uses the `otlp` receiver to ect traces.
  - **processors**: Uses the `batch` processor to batch traces.



## 1. hello-app (Microservice)

- **Role:** Acts as the main application providing the service to be monitored.
- **Instrumentation:** Instrumented with the OpenTelemetry SDK to collect and export traces and metrics.
- **Functionality:** Emits telemetry data (traces, metrics) to the OpenTelemetry Collector.
- **Communication:** Sends telemetry data to the `collector` service over HTTP.

## 2. OpenTelemetry Collector

- **Role:** Aggregates and processes telemetry data from `hello-app`.
- **Functionality:**
  - Collects traces and metrics from `hello-app`.
  - Processes and exports traces to Tempo for distributed tracing.
  - Processes and exports metrics to Prometheus for monitoring.
- **Configuration:** Uses a custom configuration file to define processing pipelines and exporters.
- **Output:**

## 3. Tempo (Tracing Backend)

- **Role:** Provides a distributed tracing backend for collecting and querying trace data.
- **Functionality:**
  - Receives trace data from the OpenTelemetry Collector.
  - Stores traces for later querying and analysis.
- **Usage:**
  - Allows Grafana to query and visualize trace data.
- **Storage:** Stores trace data in a local directory for querying.



## 2. Components

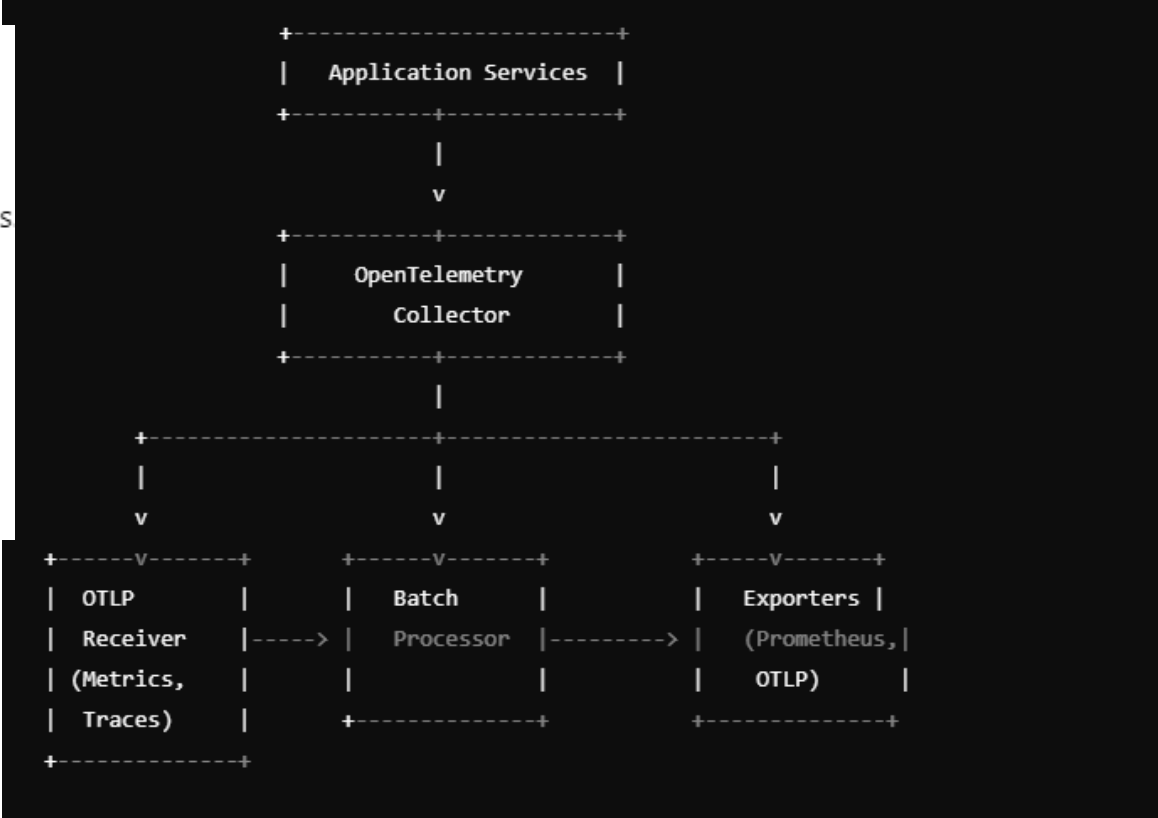
The design consists of the following key components:

- **Receivers:** Entry points for collecting telemetry data (metrics and traces) from various sources
- **Processors:** Middleware that processes, transforms, and batches telemetry data for more efficient handling.
- **Exporters:** Exit points that send telemetry data to various backends or monitoring systems.
- **Service Pipelines:** Define the flow of data from receivers through processors to exporters.

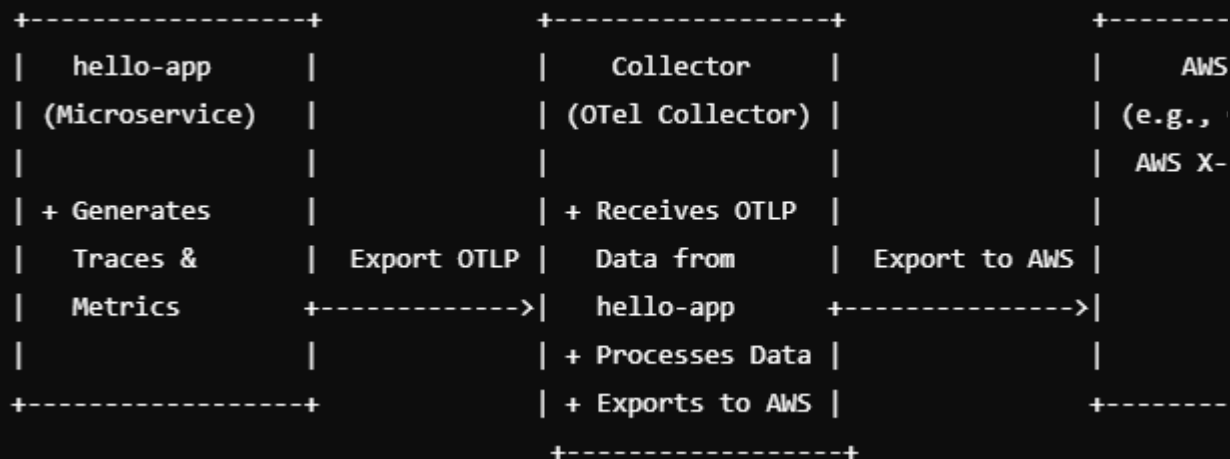
## 3. Data Flow

The data flow can be visualized as a pipeline system:

1. **Telemetry Data Collection:** Applications generate telemetry data (metrics and traces) and send it to the OpenTelemetry Collector.
2. **Reception:** The OTLP receiver listens for incoming telemetry data on a specified gRPC endpoint.
3. **Processing:** Data passes through a batch processor, which groups data into batches to optimize processing and exporting.
4. **Exporting:** Processed data is sent to configured backends:
  - Metrics are exported to a Prometheus-compatible endpoint for scraping and visualization.
  - Traces are sent to an OTLP endpoint, such as Tempo, for distributed tracing.







```
hello-app:
  build: .
  image: hello-app:latest
  container_name: hello-app
  hostname: hello-app
  depends_on:
    - collector
  ports:
    - "8888:8888"
  environment:
    - OTEL_TRACES_EXPORTER=otlp
    - OTEL_METRICS_EXPORTER=otlp
    - OTEL_EXPORTER_OTLP_ENDPOINT=http://collector:5555
    - OTEL_TRACES_SAMPLER=always_on
    - OTEL_IMR_EXPORT_INTERVAL=5000
    - OTEL_METRIC_EXPORT_INTERVAL=5000
    - OTEL_RESOURCE_ATTRIBUTES=service.name=hello-app,service.version=1.0,deployment.env=prod
```

[Copy code](#)

```
collector:
  image: public.ecr.aws/aws-observability/aws-otel-collector:latest
  container_name: collector
  hostname: collector
  command: ["--config=/etc/collector-config.yaml"]
  environment:
    - AWS_PROFILE=default
    - AWS_ACCESS_KEY_ID=AKIAQ3EGTEG7IJ6OADTH
    - AWS_SECRET_ACCESS_KEY=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
    - AWS_REGION=us-east-1
  volumes:
    - ./collector-config-aws.yaml:/etc/collector-config.yaml
```

- **environment:** Environment variables used by `hello-app`:
  - `OTEL_TRACES_EXPORTER=otlp`: Sets the OpenTelemetry traces exporter to OTLP (OpenTelemetry Protocol).
  - `OTEL_METRICS_EXPORTER=otlp`: Sets the metrics exporter to OTLP.
  - `OTEL_EXPORTER_OTLP_ENDPOINT=http://collector:5555`: Sets the endpoint for the OTLP exporter to the `collector` service running on port `5555`.
  - `OTEL_TRACES_SAMPLER=always_on`: Configures the sampler to always collect traces.
  - `OTEL_IMR_EXPORT_INTERVAL=5000`: Sets the interval for exporting interval metrics (in milliseconds).
  - `OTEL_METRIC_EXPORT_INTERVAL=5000`: Sets the interval for exporting metrics (in milliseconds).
  - `OTEL_RESOURCE_ATTRIBUTES`: Specifies resource attributes for the application, such as service name, version, and deployment environment.