


Scaling Web apps to 1,000,000 users 1/3



STARTING SIMPLE

The application started as a simple monolith. One database server and one application server. This worked well until it didn't (we started getting more users).




@MilanJovanović



REPOST

Scaling out: Multiple API instances

2/3



SCALING OUT

The next evolution was horizontally scaling the Web API to handle more requests. The database instance was scaled vertically (scale up).




@MilanJovanović



REPOST

Caching: A last resort

3/3



CACHING

The last evolution was adding a robust caching layer, which used Redis. We also added a database read replica (not in the diagram) to help for heavy reporting queries.



@MilanJovanović



REPOST